



Interpreted C++: Is that a thing?


A journey through LLVM/clang-based C++ JITting

Javier López-Gómez for the ROOT team

using `std::cpp`, 2024-04-25



Javier Lopez-Gomez

 jalopezg-git

 <https://jalopezg.dev/>

 javier.lopez.gomez AT proton.me

Summary of the last 5+ years (compilers-wise)...

- 2017–2020: PhD in Computer Science and Technology (ARCOS-UC3M)
 - Prototype implementation of C++ contracts (clang)
 - Research internship at CERN in 2019: Definition shadowing in cling
- 2020–2023: Senior Fellow (SFT, CERN)
 - More cling – but also RNTuple and general contributions to the ROOT project
- **2024–(currently): Senior Compiler Engineer (Zimperium, Inc.)**
 - Software obfuscation that operates directly AArch64 binaries



- 1 Introduction
- 2 Foundations
- 3 The cling C++ interpreter
- 4 Closing words

Introduction



You can do

```
[cling]$ template <typename T>
  T f(T a, T b) {
    return a + b;
  }
[cling]$ f(42, 6)
(int) 48
```

And then

```
[cling]$ std::string S{"Hello,"};
[cling]$ f(S, std::string{"_world!"})
(std::basic_string<char, std::char_traits<char>, std::allocator<char> >) "Hello,_world!"
```

But also the abomination below

```
[cling]$ std::vector<int> f{1, 2, 3};
[cling]$ f
(std::vector<int> &) { 1, 2, 3 }
```



You can do

```
[cling]$ template <typename T>
  T f(T a, T b) {
    return a + b;
  }
[cling]$ f(42, 6)
(int) 48
```

And then

```
[cling]$ std::string S{"Hello,"};
[cling]$ f(S, std::string{"_world!"})
(std::basic_string<char, std::char_traits<char>, std::allocator<char> >) "Hello,_world!"
```

But also the abomination below

```
[cling]$ std::vector<int> f{1, 2, 3};
[cling]$ f
(std::vector<int> &) { 1, 2, 3 }
```



You can do

```
[cling]$ template <typename T>
  T f(T a, T b) {
    return a + b;
  }
[cling]$ f(42, 6)
(int) 48
```

And then

```
[cling]$ std::string S{"Hello,"};
[cling]$ f(S, std::string{"_world!"})
(std::basic_string<char, std::char_traits<char>, std::allocator<char> >) "Hello,_world!"
```

But also the abomination below

```
[cling]$ std::vector<int> f{1, 2, 3};
[cling]$ f
(std::vector<int> &) { 1, 2, 3 }
```



- Many languages already offer a REPL (Read-Eval-Print-Loop) even if not designed to be interpreted, e.g. **C#**
- It aids a lot while learning the language: try things out!
- Iterative / exploratory prototyping
- Write 'scripts' that make use of C/C++ libraries
- ...



- Cling is built on Clang and LLVM 13 (enabling support for C++20)
- CUDA support
- Allows loading an external library (`.so` / `.dll`) and get access to its symbols, e.g. call a function
- Debugging and profiling of JITed code
- Undo steps
- Protection against invalid memory accesses, e.g. dereferencing a pointer that points to unmapped memory



But also some features that one would expect from an interpreter (even if that's not ISO C++)...

- Top-level statements
- Print the result of expression evaluation
- `auto` synthesizing, i.e. `foo = 42.0;` is equivalent to `auto foo = 42.0;` (if `foo` not declared before) DEPRECATED
- Support for redefining entities, e.g.

```
int foo = 0;  
std::string foo{"Hello!"};
```

Foundations



- LLVM and clang to the rescue!



- LLVM gives all the infrastructure required to build a compiler
 - Basic data structures, e.g. `llvm::SmallVector`, or `llvm::Twine`
 - An intermediate representation (LLVM IR)
 - Lowering to machine code for many targets: x86_64, ARM7, AArch64, RISC-V, etc.
 - Machine-dependent and machine-independent optimization passes
 - Generation of debug information



- LLVM and clang to the rescue!

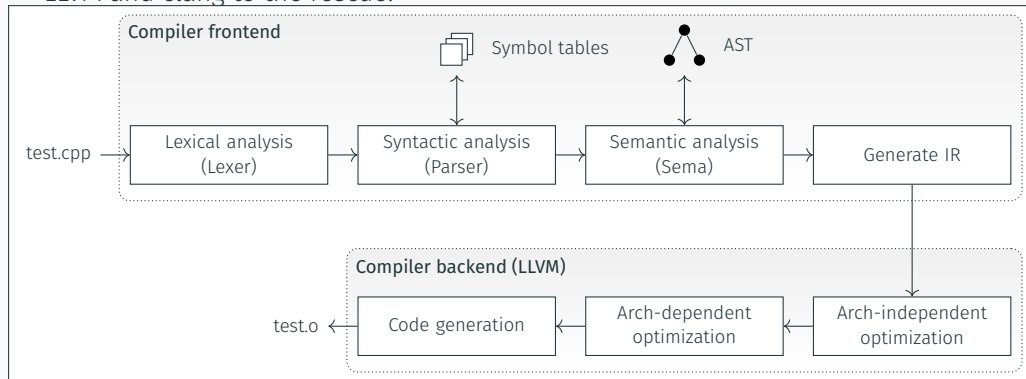


- LLVM gives all the infrastructure required to build a compiler
 - Basic data structures, e.g. `llvm::SmallVector`, or `llvm::Twine`
 - An intermediate representation (LLVM IR)
 - Lowering to machine code for many targets: x86_64, ARM7, AArch64, RISC-V, etc.
 - Machine-dependent and machine-independent optimization passes
 - Generation of debug information

And more! See <https://llvm.org/>.



- LLVM and clang to the rescue!



And more! See <https://llvm.org/docs/GettingStarted.html>



- LLVM IR may have 3 different representations: **in-memory** structures, assembly-like **plain-text**, or **serialized** form
- Let's play a bit to build the LLVM IR for a simple function!



```
LLVMContext C;
auto builder = std::make_unique<IRBuilder<>>(C);
auto M = std::make_unique<Module>("main", C);

auto i32 = builder->getInt32Ty();
auto funcTy = FunctionType::get(i32, {i32, i32}, /*isVarArg=*/false);
auto func = Function::Create(funcTy, GlobalValue::LinkageTypes::ExternalLinkage, "func", *M);
auto BB = BasicBlock::Create(C, "entry", func);
builder->SetInsertPoint(BB);
auto addVal = builder->CreateAdd(func->getArg(0), func->getArg(1));
builder->CreateRet(addVal);
M->print(errs(), /*AAW=*/nullptr);
```




```

LLVMContext C;
auto builder = std::make_unique<IRBuilder<>>(C);
auto M = std::make_unique<Module>("main", C);

auto i32 = builder->getInt32Ty();
auto funcTy = FunctionType::get(i32, {i32, i32});
auto func = Function::Create(funcTy, GlobalValue::ExternalLinkage, M);
auto BB = BasicBlock::Create(C, "entry", func);
builder->SetInsertPoint(BB);
auto addVal = builder->CreateAdd(func->getArg(0), func->getArg(1));
builder->CreateRet(addVal);
M->print(errs(), /*AAW=*/nullptr);

```

```

; ModuleID = 'main'
source_filename = "main"

define i32 @func(i32 %0, i32 %1) {
entry:
    %2 = add i32 %0, %1
    ret i32 %2
}

```



LLVM can also JIT IR to current target's machine code...¹

```
int main(int argc, char *argv[]) {
    using namespace llvm;
    InitializeNativeTarget();
    InitializeNativeTargetAsmPrinter();

    /* CREATE IR AS IN PREVIOUS SLIDE */

    auto EE = EngineBuilder(std::move(M)).setEngineKind(llvm::EngineKind::JIT).create();

    using FuncPtr_t = uint32_t (*)(uint32_t, uint32_t);
    auto pFunc = (FuncPtr_t)EE->getFunctionAddress("func");
    auto ret = pFunc(42, 7);
    errs() << "\nfunc() returned " << ret << "\n";

    return 0;
}
```

¹Full example: https://github.com/jalopezg-git/slides-using_stdcpp_2014/blob/master/code/llvm-ir.cpp



LLVM can also JIT IR to current target's machine code...¹

```
int main(int argc, char *argv[]) {  
    using namespace llvm;  
    InitializeNativeTarget();  
    InitializeNativeTargetAsmPrinter();  
  
    /* CREATE IR AS IN PREVIOUS SLIDE */  
  
    auto EE = EngineBuilder(std::move(M)).setEngineKind(llvm::EngineKind::JIT).create();  
  
    using FuncPtr_t = uint32_t (*)(uint32_t, uint32_t);  
    auto pFunc = (FuncPtr_t)EE->getFunctionAddress("func");  
    auto ret = pFunc(42, 7);  
    errs() << "\nfunc() returned " << ret << "\n";  
  
    return 0;  
}
```

func() returned 49

¹Full example: https://github.com/jalopezg-git/slides-using_stdcpp_2014/blob/master/code/llvm-ir.cpp



- Clang is basically a frontend that parses C / C++ / ObjectiveC and generates LLVM IR, i.e.
 - It does lexical / grammatical / semantic analysis on the source code + builds an AST
 - LLVM takes over from there
- E.g., the simple code...

```
extern int puts(const char *s);

int main(void) {
    puts("Hello, \uworld!");
    return 0;
}
```

²Get this with `clang -c -Xclang -ast-dump -o /dev/stdout input.c`

³Get this with `clang -S -emit-llvm -o /dev/stdout input.c`



```
extern int puts(const char *s);

int main(void) {
    puts("Hello, world!");
    return 0;
}
```

Has AST representation²

```
-FunctionDecl 0x563c61b0bda0 </tmp/simple.c:1:1, col:30> col:12 used puts 'int (const char *)' extern
|-ParmVarDecl 0x563c61b0bcd0 <col:17, col:29> col:29 s 'const char *'
|-FunctionDecl 0x563c61b0bf60 <line:3:1, line:6:1> line:3:5 main 'int (void)'
|-CompoundStmt 0x563c61b0c188 <col:16, line:6:1>
| |-CallExpr 0x563c61b0c100 <line:4:5, col:25> 'int'
| | |-ImplicitCastExpr 0x563c61b0c0e8 <col:5> 'int (*)(const char *)' <FunctionToPointerDecay>
| | | |-DeclRefExpr 0x563c61b0c038 <col:5> 'int (const char *)' Function 0x563c61b0bda0 'puts' 'int (const char *)'
| | | |-ImplicitCastExpr 0x563c61b0c140 <col:10> 'const char *' <NoOp>
| | | |-ImplicitCastExpr 0x563c61b0c128 <col:10> 'char *' <ArrayToPointerDecay>
| | | | |-StringLiteral 0x563c61b0c098 <col:10> 'char[14]' lvalue "Hello, world!"
| | -ReturnStmt 0x563c61b0c178 <line:5:5, col:12>
| -IntegerLiteral 0x563c61b0c158 <col:12> 'int' 0
```

²Get this with `clang -c -Xclang -ast-dump -o /dev/stdout input.c`

³Get this with `clang -S -emit-llvm -o /dev/stdout input.c`



```
extern int puts(const char *s);

int main(void) {
    puts("Hello, world!");
    return 0;
}
```

And LLVM IR representation³

```
; ModuleID = '/tmp/simple.c'
source_filename = "/tmp/simple.c"
target datalayout = "e-m:e-p270:32:32-p271:32:32-p272:64:64-i64:64-f80:128-n8:16:32:64-S128"
target triple = "x86_64-pc-linux-gnu"
```

```
@.str = private unnamed_addr constant [14 x i8] c"Hello, world!\00", align 1
```

```
; Function Attrs: nofree nounwind sspstrong uwtable
define dso_local i32 @main() local_unnamed_addr #0 {
    %1 = tail call i32 @puts(ptr noundef nonnull dereferenceable(1) @.str)
    ret i32 0
}
```

```
; Function Attrs: nofree nounwind
declare noundef i32 @puts(ptr nocapture noundef readonly) local_unnamed_addr #1
```

²Get this with `clang -c -Xclang -ast-dump -o /dev/stdout input.c`

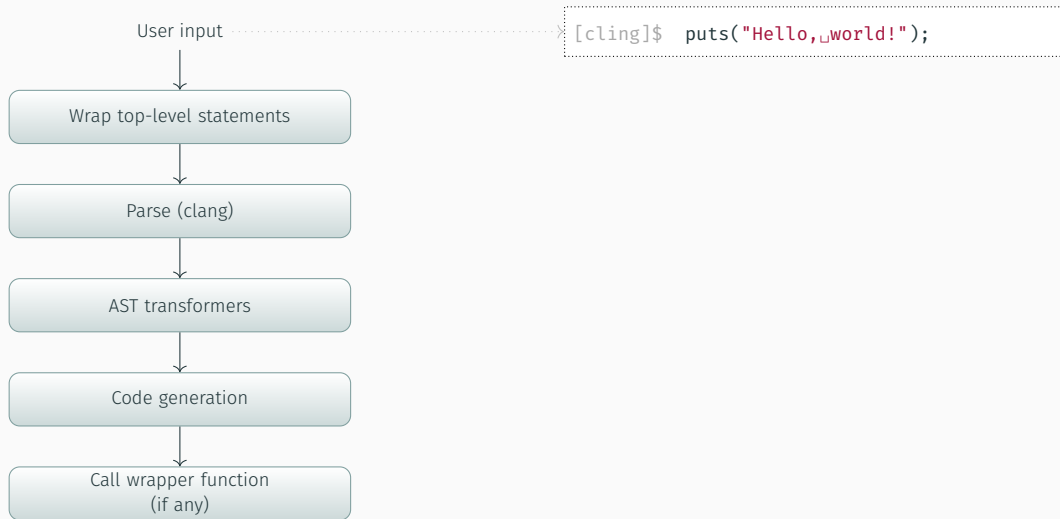
³Get this with `clang -S -emit-llvm -o /dev/stdout input.c`



- And it offers libTooling, libclang-cpp, and libclang!
- Meaning we can mostly reuse this⁴ and only write a layer on top that does “impedance matching” between ISO and interpreted C++

⁴Modulo some patches required to clang.

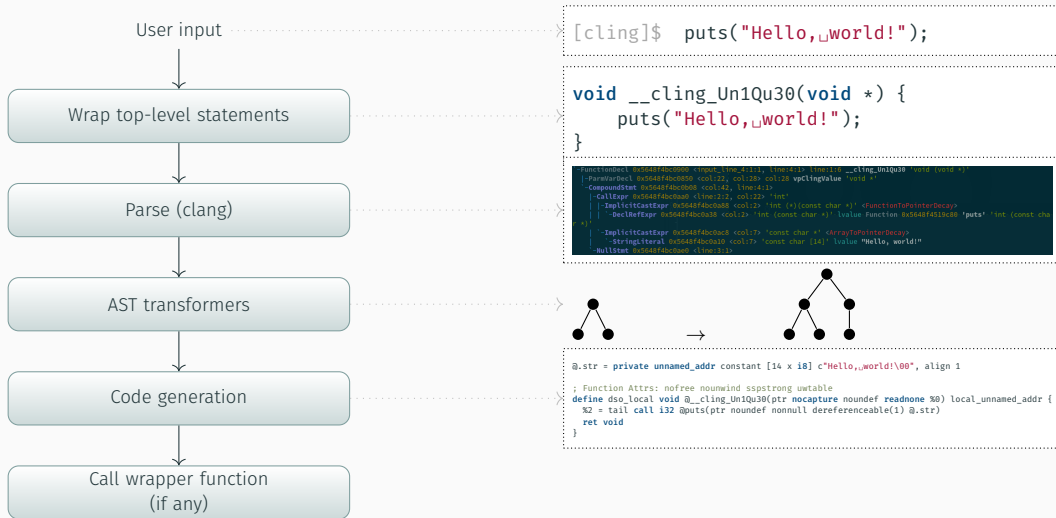
The cling C++ interpreter

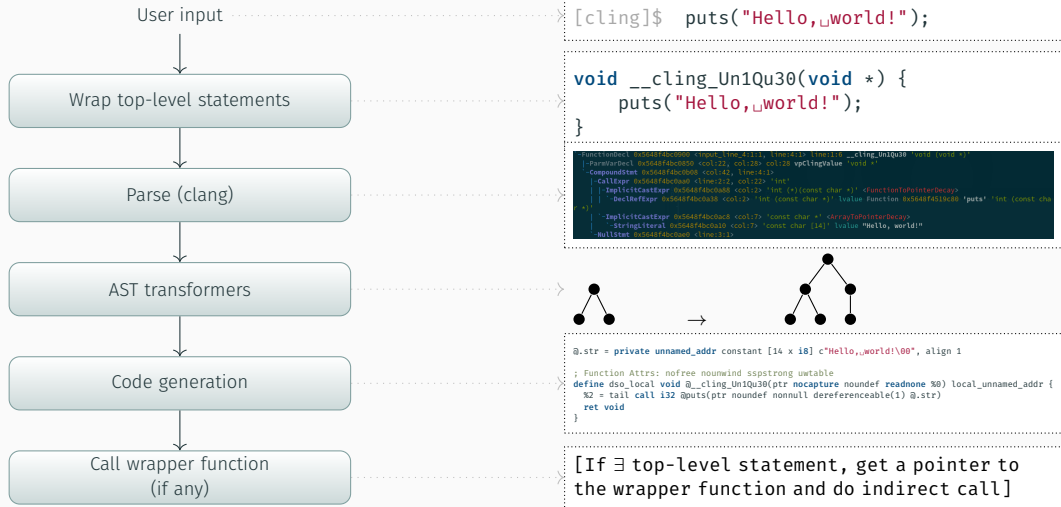


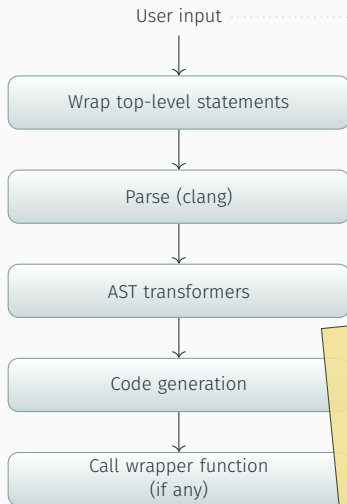












```
[cling]$ puts("Hello, world!");
```

```
void __cling_Un1Qu30(void *) {  
    puts("Hello, world!");  
}
```

```
FunctionDecl @x5648f4bc0a0e <input_line:4:11:1, line:4:1> {  
  ParamDecl @x5648f4bc0a0e <col:22, col:28> col:28 vpClingValue 'void *'  
  CompoundStmt @x5648f4bc0a08 <col:42, line:4:1>  
    CallExpr @x5648f4bc0a08 <line:2:2, col:22> 'int'  
      ImplicitCastExpr @x5648f4bc0a08 <col:2> 'int (*) (const char *)' <FunctionToPointerDecay>  
      DeclRefExpr @x5648f4bc0a38 <col:2> 'int (const char *)' lvalue Function @x5648f4519c80 'puts' 'int (const cha  
  }  
  ImplicitCastExpr @x5648f4bc0a08 <col:7> 'const char *' <ArrayToPointerDecay>  
  StringLiteral @x5648f4bc0a10 <col:7> 'const char [14]' lvalue "Hello, world!"  
  NullStmt @x5648f4bc0a08 <line:3:3>
```

- Deferred (implicit) template instantiations must be emitted; we do that by forcing a end-of-TU event!
- CodeGen: also involves linking (external symbol resolution, etc.)



- AST is built incrementally
- **Transaction:** declarations that were parsed and emitted in a single step
 - User-provided declarations
 - Implicit template instantiations
 - Deserialized declarations from a C++ module
- And allows undoing it. That's useful, e.g. after a failed parse



- Most extensions are implemented as an AST transformer
- Currently, there is support for
 - `auto` synthesizing, e.g. `foobar = 42.0f;`
 - Protection against invalid pointer dereferencing, e.g.

```
[cling]$ *((int *)0xff00ff00) = 0;
Error in <HandleInterpreterException>: Trying to access a pointer that points to an
invalid memory address.
Execution of your code was aborted.
ROOT_prompt_6:1:2: warning: invalid memory pointer passed to a callee:
*((int *)0xff00ff00) = 0;
  ^~~~~~
```

- Shadowing of definitions

```
[cling]$ int foobar = 0;
[cling]$ std::string foobar() { return "A_string!"; }
```

- Printing / capturing the value of an expression

```
[cling]$ foobar()
(std::string) "A_string!"
```



- Most extensions are implemented as an AST transformer
- Currently, there is support for
 - `auto` synthesizing, e.g. `foobar = 42.0f;`
 - Protection against invalid pointer dereferencing, e.g.

```
[cling]$ *((int *)0xff00ff00) = 0;
Error in <HandleInterpreterException>: Trying to access a pointer that points to an
invalid memory address.
Execution of your code was aborted.
ROOT_prompt_6:1:2: warning: invalid memory pointer passed to a callee:
*((int *)0xff00ff00) = 0;
  ^~~~~~
```

- Shadowing of definitions

```
[cling]$ int foobar = 0;
[cling]$ std::string foobar() { return "A_string!"; }
```

- Printing / capturing the value of an expression

```
[cling]$ foobar()
(std::string) "A_string!"
```



- Most extensions are implemented as an AST transformer
- Currently, there is support for
 - `auto` synthesizing, e.g. `foobar = 42.0f;`
 - Protection against invalid pointer dereferencing, e.g.

```
[cling]$ *((int *)0xff00ff00) = 0;
Error in <HandleInterpreterException>: Trying to access a pointer that points to an
invalid memory address.
Execution of your code was aborted.
ROOT_prompt_6:1:2: warning: invalid memory pointer passed to a callee:
*((int *)0xff00ff00) = 0;
  ^~~~~~
```

- Shadowing of definitions

```
[cling]$ int foobar = 0;
[cling]$ std::string foobar() { return "A_string!"; }
```

- Printing / capturing the value of an expression

```
[cling]$ foobar()
(std::string) "A_string!"
```



- Most extensions are implemented as an AST transformer
- Currently, there is support for
 - `auto` synthesizing, e.g. `foobar = 42.0f;`
 - Protection against invalid pointer dereferencing, e.g.

```
[cling]$ *((int *)0xff00ff00) = 0;
Error in <HandleInterpreterException>: Trying to access a pointer that points to an
    invalid memory address.
Execution of your code was aborted.
ROOT_prompt_6:1:2: warning: invalid memory pointer passed to a callee:
*((int *)0xff00ff00) = 0;
  ^~~~~~
```

- Shadowing of definitions

```
[cling]$ int foobar = 0;
[cling]$ std::string foobar() { return "A_string!"; }
```

- Printing / capturing the value of an expression

```
[cling]$ foobar()
(std::string) "A_string!"
```



Cling also allows debugging JITed code and offers integration with Linux's **perf**, e.g.

- A breakpoint on interpreted code can be set and step-into after each statement
- It can generate a symbol file for **perf** – Can be used together with Flamegraph⁵!

⁵Flamegraph: <https://github.com/brendangregg/FlameGraph>



- Cling proved to perform okay in the context of the larger ROOT project at CERN
- Let's upstream the foundations of it back to the LLVM community so that
 - The whole community can benefit from it
 - Maintenance is easier in the long term
- `clang-repl`: already in recent versions of LLVM — Thanks, Vassil!
- Slightly different to the design of cling, e.g. modeling of top-level statements is much more robust

Closing words



Key ideas to take home

- Cling enables incremental C++ compilation and JITting
 - It would not be possible without the solid framework provided by LLVM and clang
- Convenient integration with Jupyter notebook via **xeus-cling**
- Try it!

For the curious

- cppy / libinterop provide interoperability with other languages
 - Enabling crazy things such as injecting the C++ definition of a type **T** and creating an object of type **T** from Python
 - Or even crazier: on-the-fly template instantiation, e.g. a **std::vector<T>**^{ab}

^a<https://cpypy.readthedocs.io/>

^b<https://compiler-research.org/libinterop/>



<https://github.com/root-project/cling/>

- If you have 'llvm' installed locally, try `clang-repl`



Thank you!

Backup



Cling also allows debugging JITed code and offers integration with Linux's **perf**, e.g.

- A breakpoint on interpreted code can be set and step-into after each statement

```
$ export CLING_DEBUG=1
$ gdb --args cling /tmp/simple.C
(gdb) break simple
(gdb) r
Starting program: cling /tmp/simple.C


Breakpoint 1, simple () at /tmp/simple.C:4
4      std::cout << "Hello, world!" << std::endl;
(gdb) q
```

- It can generate a symbol file for **perf** – Can be used together with Flamegraph⁶!

```
$ export CLING_PROFILE=1
$ perf record -g -e cycles -- cling /tmp/simple.C
```

⁶Flamegraph: <https://github.com/brendangregg/FlameGraph>

Debugging / Profiling of JIT'ed code: flamegraph



Interpreted C++: Is that a thing?

using std::cpp, 2024-04-25